



## Document Processing Architecture

- **Input Rate:** 5M docs/day = ~3,470 docs/minute (1440 batches/day)
- **Batch Processing Pattern:** Documents delivered every 60 seconds by internal process
- **Peak Batch Size:** ~3,470 documents per batch (with potential variation  $\pm 20\%$ )
- **Processing Time per Document:** Estimated 2-5 seconds (parallel ML inference)
- **Lambda Concurrency per Batch:** ~3,470 concurrent executions (auto-scaling)
- **Batch Processing SLA:** Complete processing within 5 minutes of batch receipt
- **Cost Effective:** 96% cheaper than PostgreSQL/Cassandra (~\$120 vs \$3100/month)
- **Unlimited Scale:** No capacity planning needed, grows with your data
- **Analytics Ready:** Native integration with Athena, Glue, EMR for data science
- **Durability:** 99.999999999% (11 9's) durability, automatic replication
- **Multiple Storage Classes:** Intelligent tiering for cost optimization